

Introduction to Deep Learning

7- Architectures (Part 2): Recurrent Neural Networks (RNNs)

Prof. Monir EL ANNAS

Networks for sequential data

Sequential Data:

- Most of data are sequential: Speech, Text, Image, ...
- There are 3 problem types: Predict the next step, classify, and generate a sequence

Problems with feedforward neural networks:


- They assume that each data point is independent and non-sequential.
- They lack the capability to retain any memory of previous inputs, treating each new input in isolation.

Deep Learnings for Sequential Data:

- Convolutional Neural Networks : Try to find local features from a sequence
- Recurrent Neural Networks : Try to capture the feature of the past

Recurrent neural networks

Recurrent neural networks: Structure



$$h_t = f(Ux_t + Wh_{t-1})$$

$$o_t = g(Vh_t)$$


X_t : t^{th} input (sequential: X_1, X_2, \dots, X_t)

h_t : t^{th} hidden state

f : activation function

U, V, W : Network parameters



g : Activation function of the output layer



Recurrent neural networks

Recurrent neural networks: Structure

- **Sequential Data Processing:** Each network pass handles one time step of input, influencing the next.
- **Shared Parameters:** RNNs reuse the same weights at each step.
- **Backpropagation Through Time:** Unfolding the network through each time step, allow gradient calculations over sequences.



$$h_t = \tanh(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$


Recurrent neural networks

Recurrent neural networks: Different combinations

- Different applications will give rise to different ways in which we use RNNs

Feedforward


Neural Network



one to one


Image
classification

Recurrent Neural Network




one to many

Image
captioning




many to one

Video
classification



many to many

Machine
translation




many to many

Per-frame video
classification

Recurrent neural networks

Recurrent Neural Networks: Multi-layer RNN

- Several layers of RNN units (cells) can be stacked vertically. Each layer processes the sequence output from the previous layer.
- The layers usually use different weight matrices.



Recurrent neural networks

Recurrent Neural Networks: PyTorch

```
CLASS torch.nn.RNN(self, input_size, hidden_size, num_layers=1, nonlinearity='tanh',  
                  bias=True, batch_first=False, dropout=0.0, bidirectional=False, device=None,  
                  dtype=None) [SOURCE]
```



Apply a multi-layer Elman RNN with `tanh` or `ReLU` non-linearity to an input sequence. For each element in the input sequence, each layer computes the following function:

$$h_t = \tanh(x_t W_{ih}^T + b_{ih} + h_{t-1} W_{hh}^T + b_{hh})$$

where h_t is the hidden state at time t , x_t is the input at time t , and h_{t-1} is the hidden state of the previous layer at time $t-1$ or the initial hidden state at time 0. If `nonlinearity` is `'relu'`, then `ReLU` is used instead of `tanh`.


Parameters

- **input_size** – The number of expected features in the input x
- **hidden_size** – The number of features in the hidden state h
- **num_layers** – Number of recurrent layers. E.g., setting `num_layers=2` would mean stacking two RNNs together to form a *stacked RNN*, with the second RNN taking in outputs of the first RNN and computing the final results. Default: 1
- **nonlinearity** – The non-linearity to use. Can be either `'tanh'` or `'relu'`. Default: `'tanh'`
- **bias** – If `False`, then the layer does not use bias weights `b_ih` and `b_hh`. Default: `True`
- **batch_first** – If `True`, then the input and output tensors are provided as `(batch, seq, feature)` instead of `(seq, batch, feature)`. Note that this does not apply to hidden or cell states. See the Inputs/Outputs sections below for details. Default: `False`
- **dropout** – If non-zero, introduces a *Dropout* layer on the outputs of each RNN layer except the last layer, with dropout probability equal to `dropout`. Default: 0
- **bidirectional** – If `True`, becomes a *bidirectional RNN*. Default: `False`

Recurrent neural networks

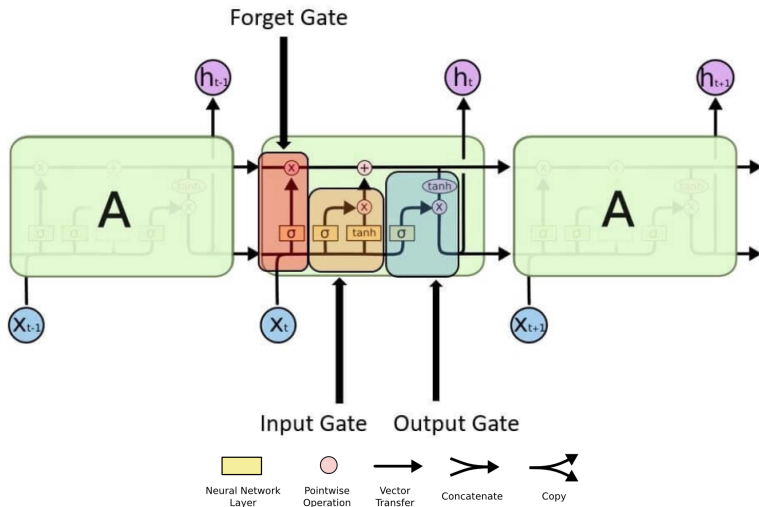
Recurrent Neural Networks: Problem

- Long-term dependency problem :
 - Must remember all states at any given time
 - Computationally expensive
 - Only store states within a time window
- Sensitive to changes in their parameters
- Vanishing Gradient
- Exploding Gradient




RNNs variants

Long Short-Term Memory (LSTM)




RNNs variants

Long Short-Term Memory (LSTM)



RNNs variants


Long Short-Term Memory (LSTM)



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

RNNs variants

Long Short-Term Memory (LSTM)




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

RNNs variants


Long Short-Term Memory (LSTM)



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

RNNs variants

Long Short-Term Memory (LSTM)




$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

RNNs variants

Gated Recurrent Unit (GRU)



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$